



Red Hat

Ansible Automation
Platform

Getting Started with Ansible

What is the Red Hat Ansible Automation Platform?

The Ansible project is an open source community sponsored by Red Hat. It's also a simple automation language that perfectly describes IT application environments in Ansible Playbooks.

Ansible Engine is a supported product built from the Ansible community project.


Ansible Tower is an enterprise framework for controlling, securing, managing and extending your Ansible automation (community or engine) with a UI and RESTful API.

Why Ansible?



Simple

Human readable automation
No special coding skills needed
Tasks executed in order
Usable by every team
Get productive quickly



Powerful

App deployment
Configuration management
Workflow orchestration
Network automation
Orchestrate the app lifecycle



Agentless

Agentless architecture
Uses OpenSSH & WinRM
No agents to exploit or update
Get started immediately
More efficient & more secure

With Ansible you can automate:

CROSS PLATFORM – Linux, Windows, UNIX

Agentless support for all major OS variants, physical, virtual, cloud and network

HUMAN READABLE – YAML

Perfectly describe and document every aspect of your application environment

PERFECT DESCRIPTION OF APPLICATION

Every change can be made by playbooks, ensuring everyone is on the same page

VERSION CONTROLLED

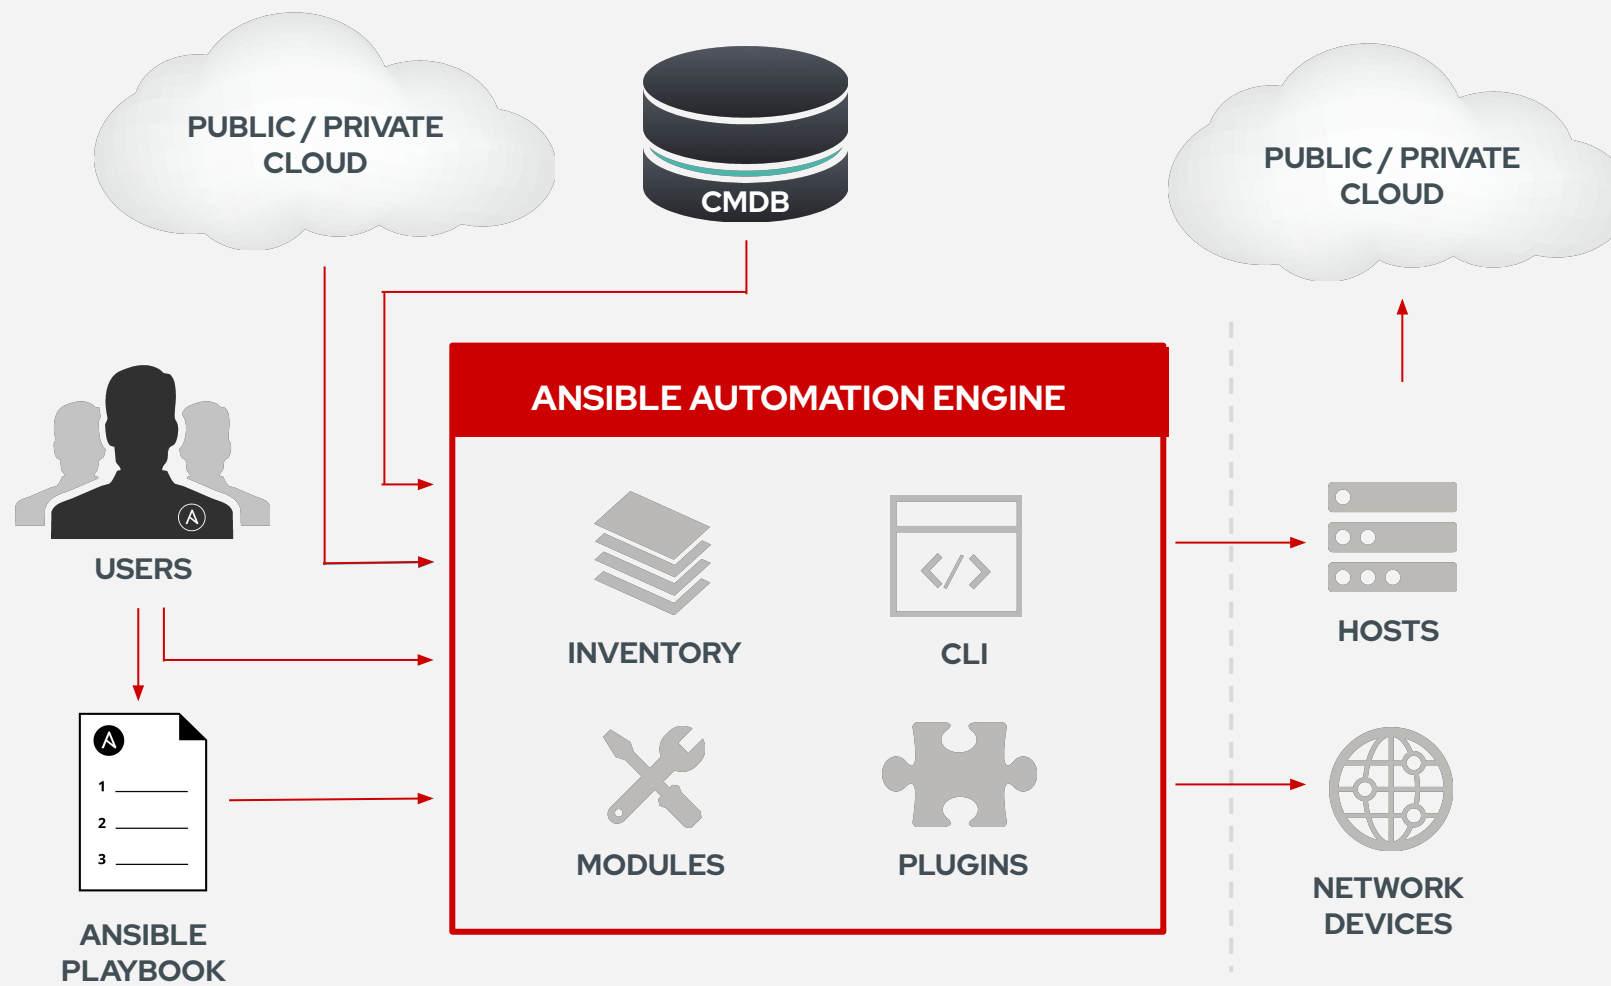
Playbooks are plain-text. Treat them like code in your existing version control.

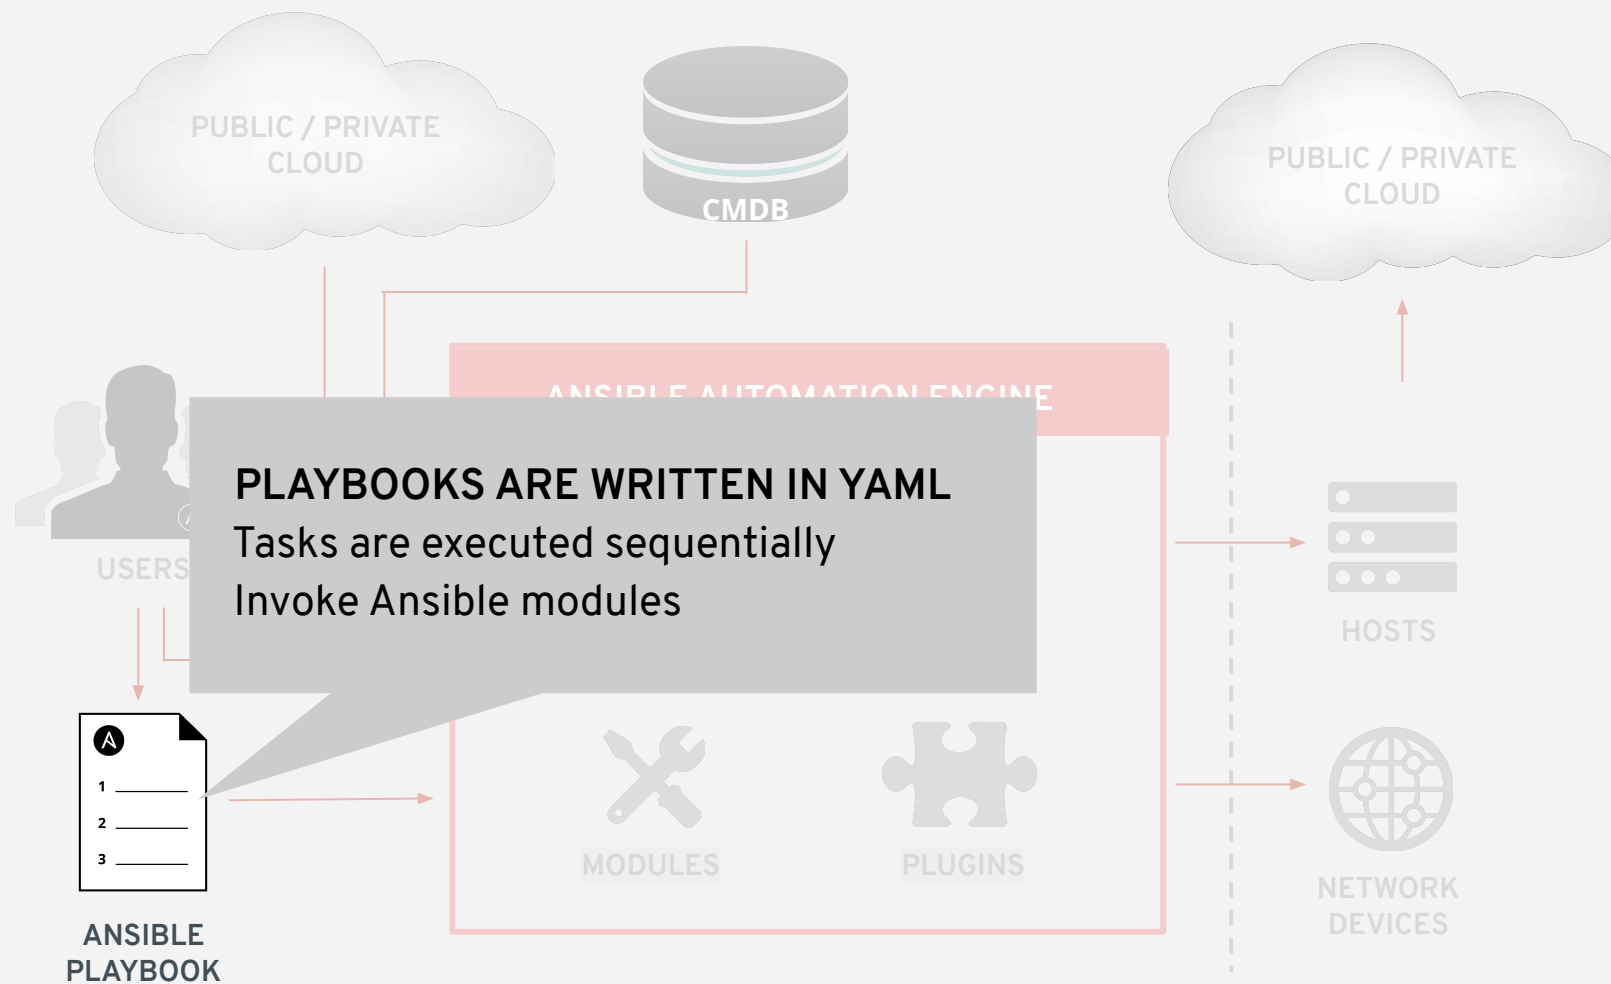
DYNAMIC INVENTORIES

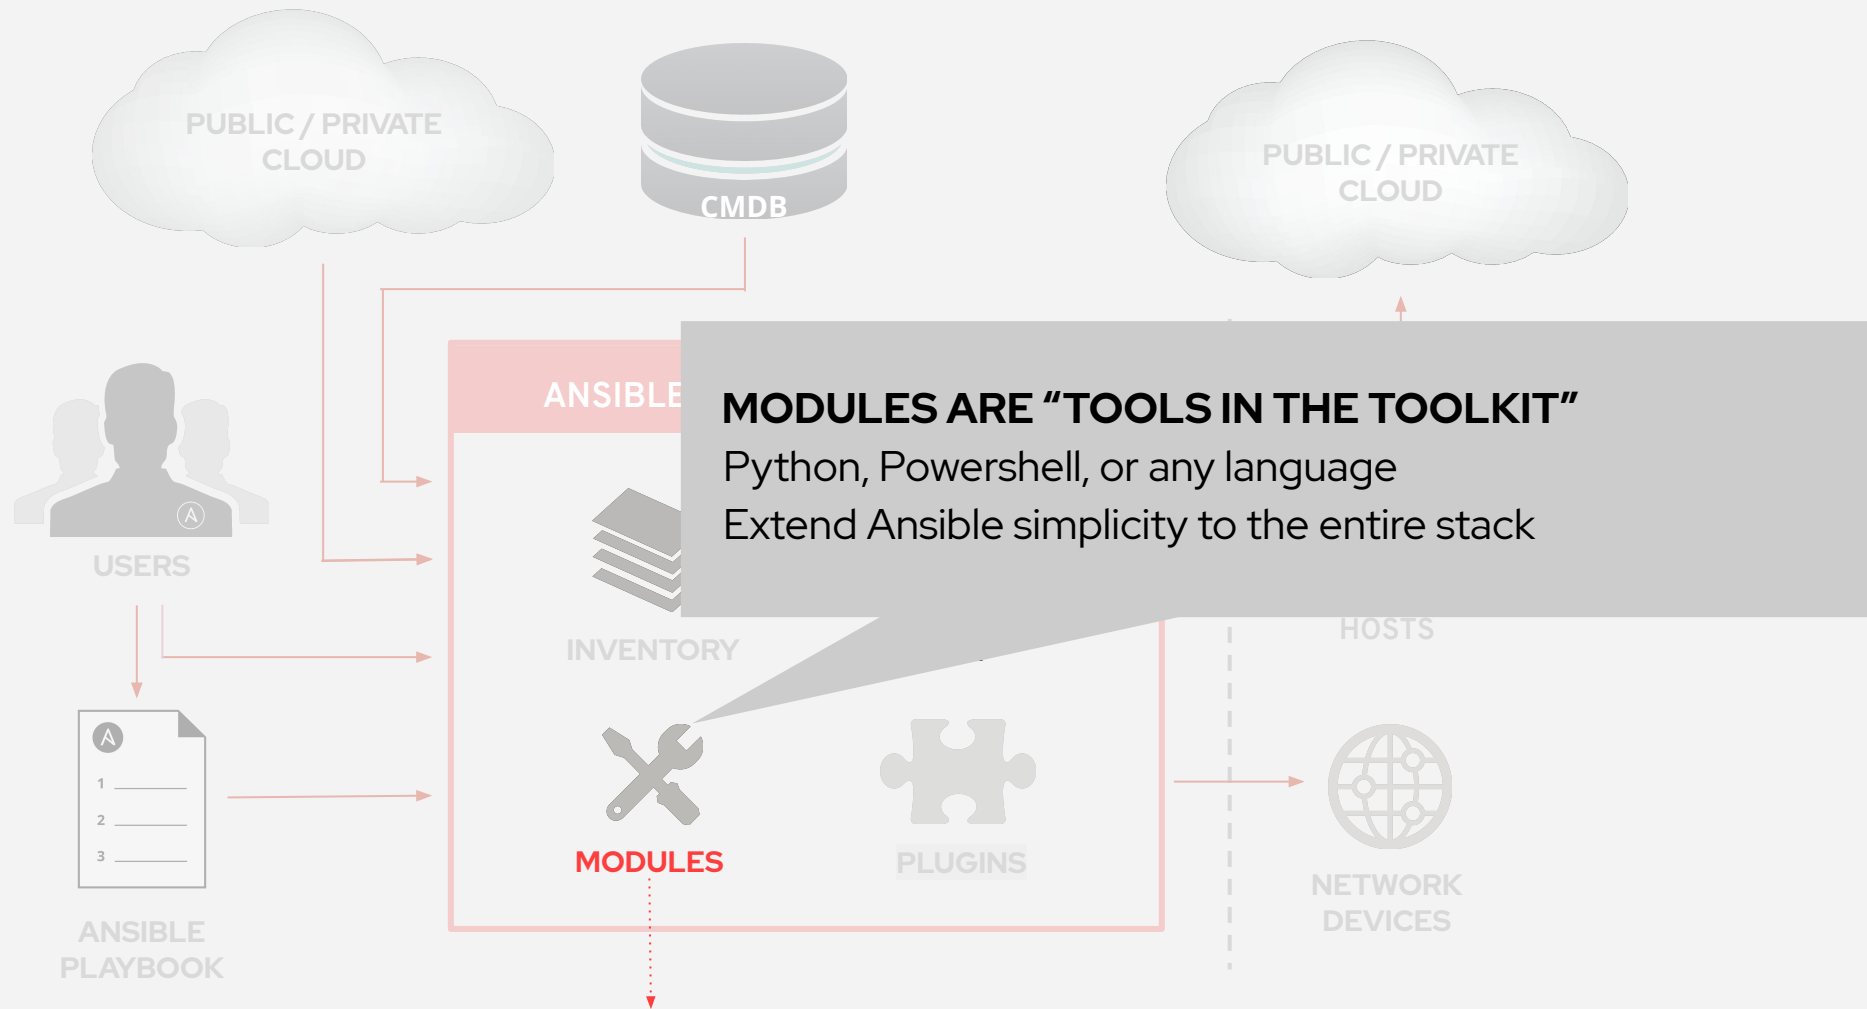
Capture all the servers 100% of the time, regardless of infrastructure, location, etc.

ORCHESTRATION THAT PLAYS WELL WITH OTHERS – HP SA, Puppet, Jenkins, RHNSS, etc.

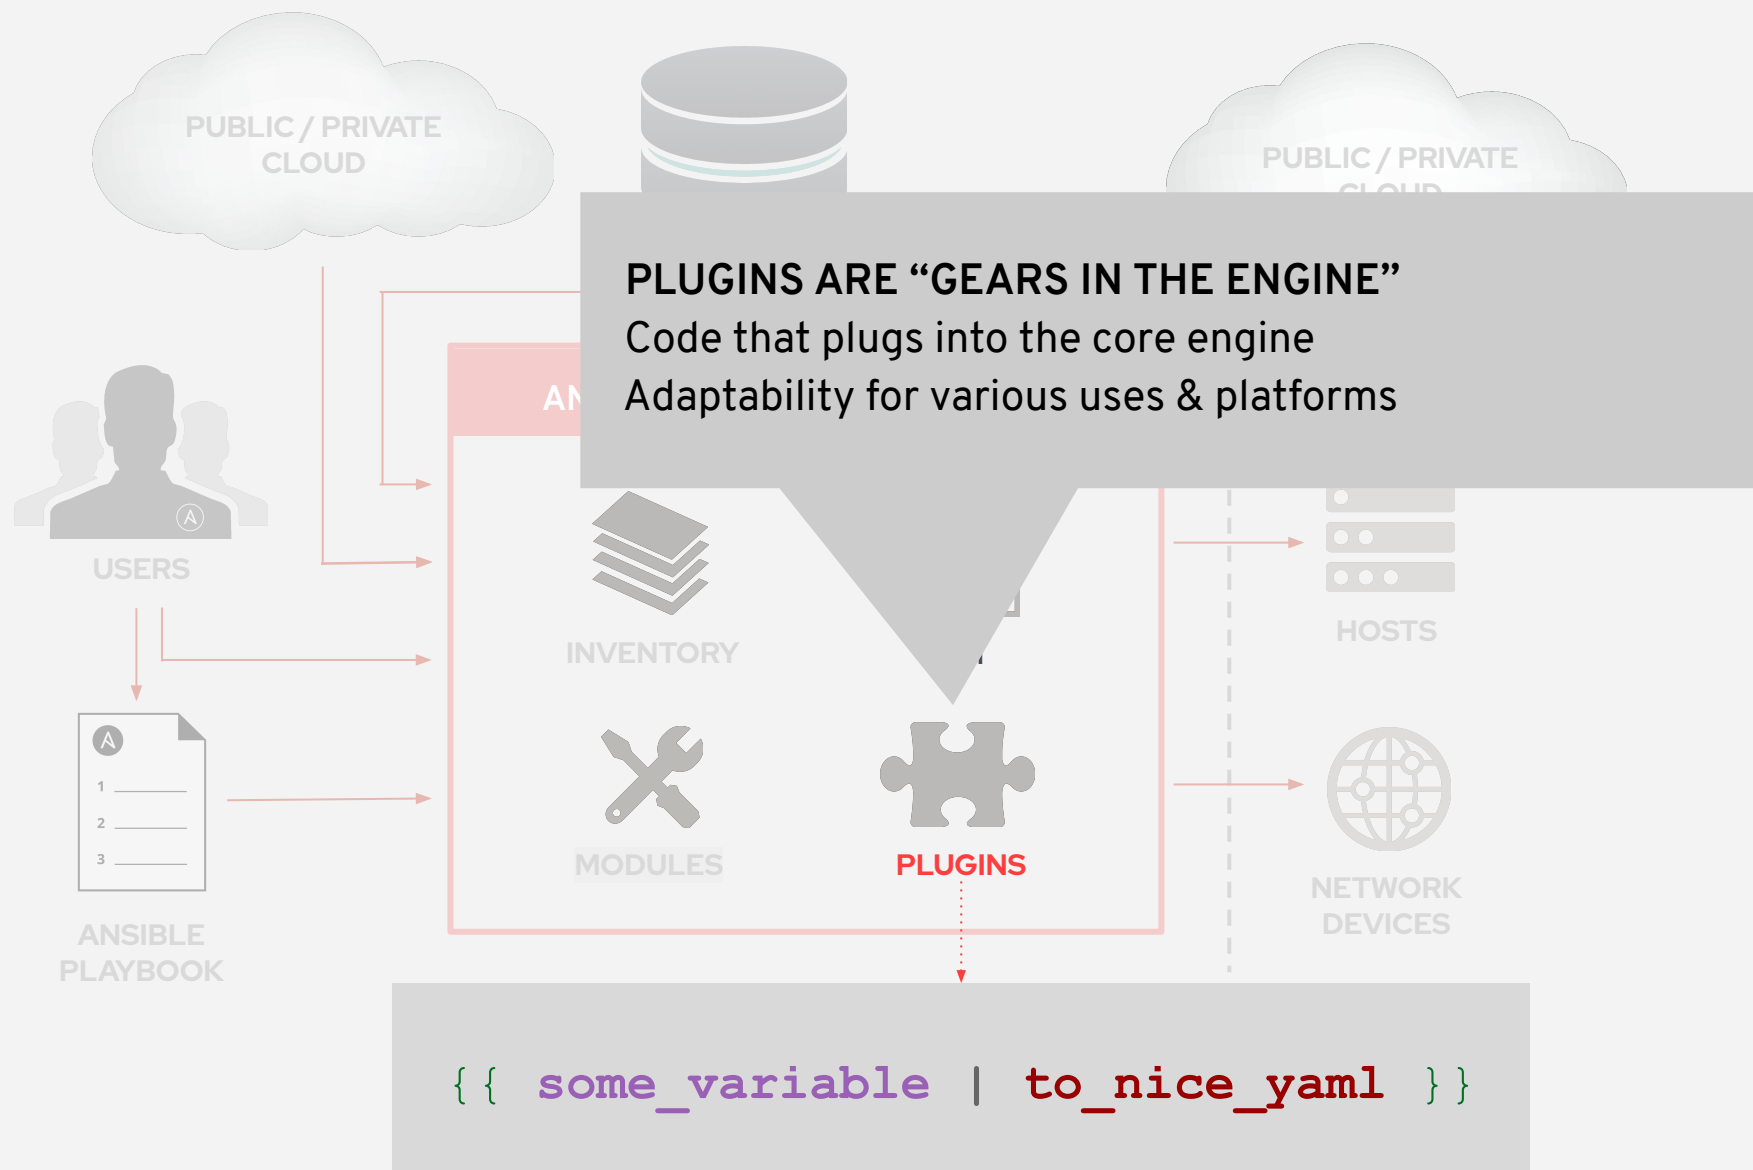
Homogenize existing environments by leveraging current toolsets and update mechanisms.

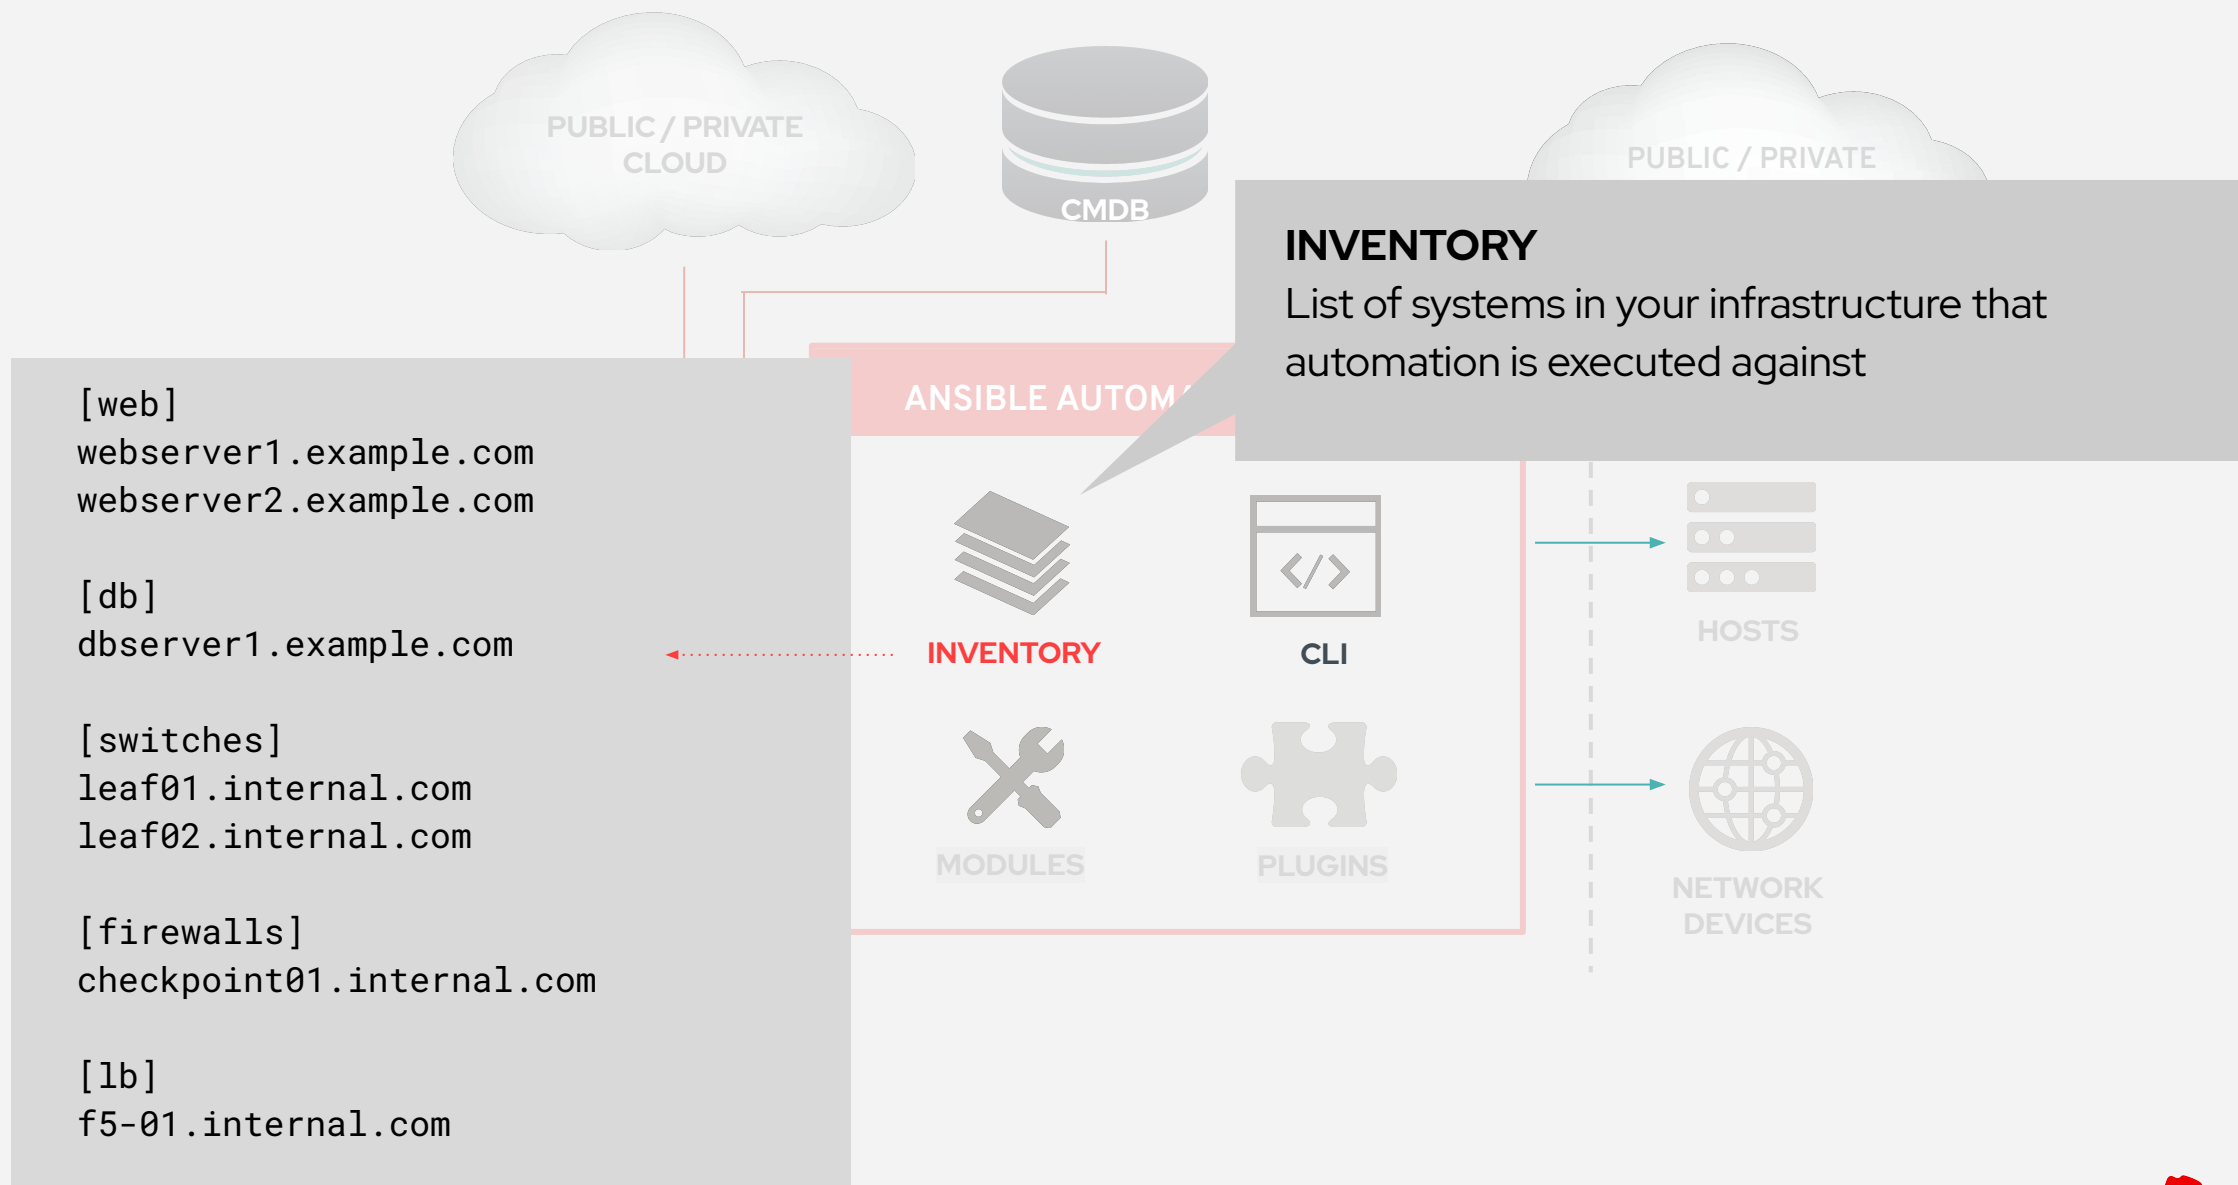


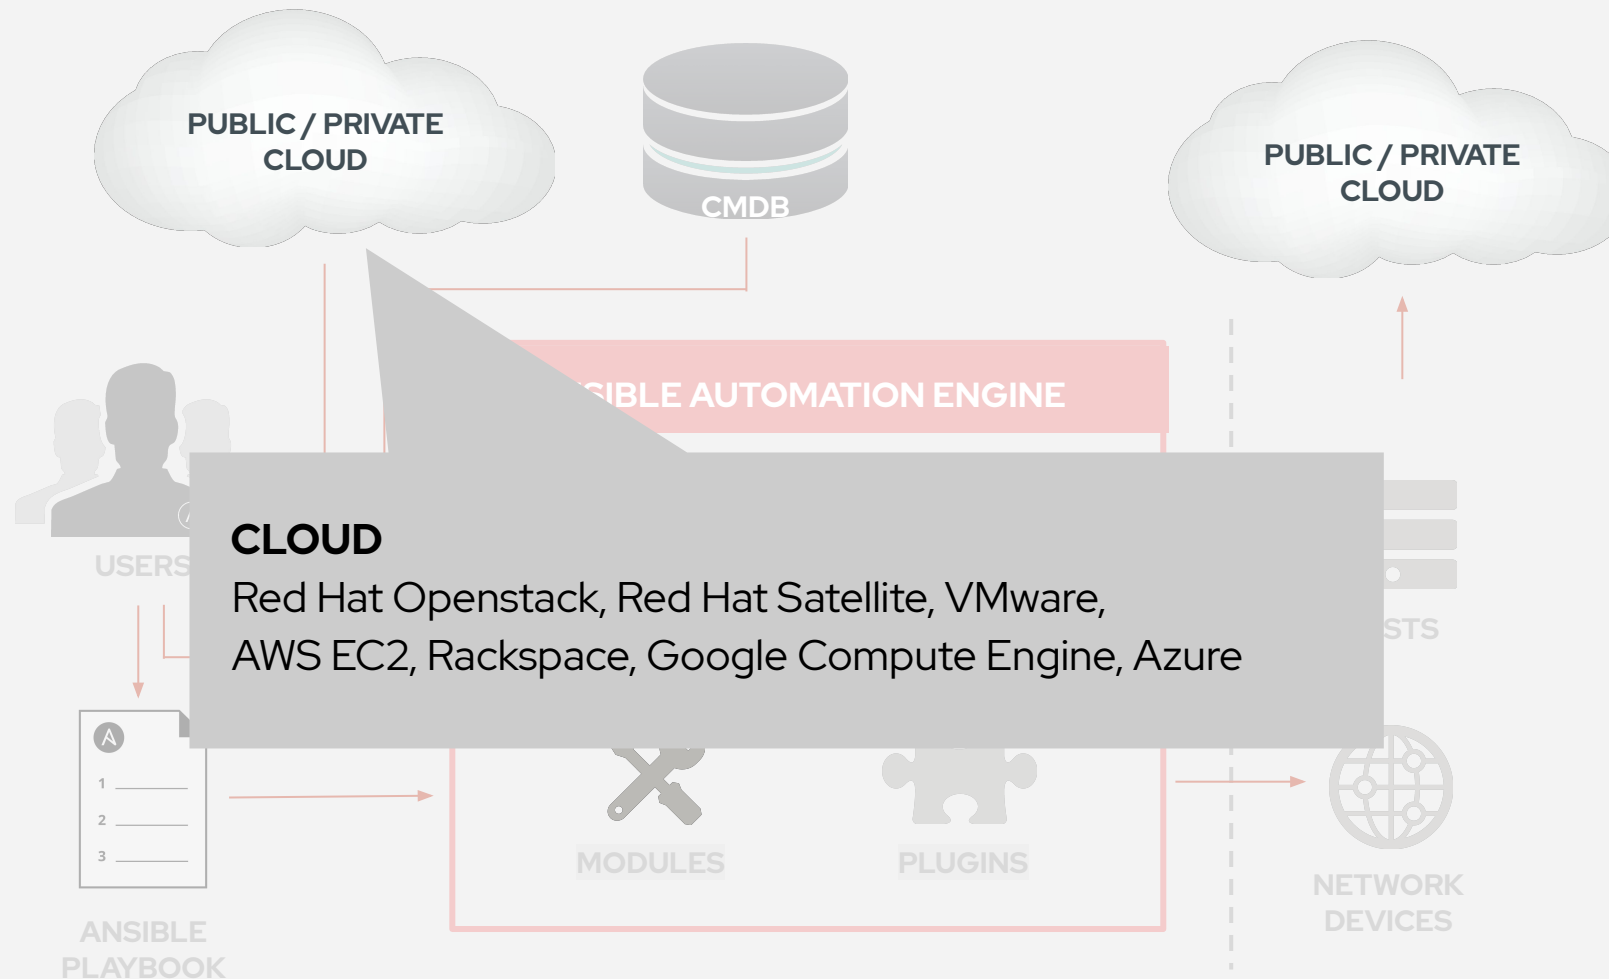


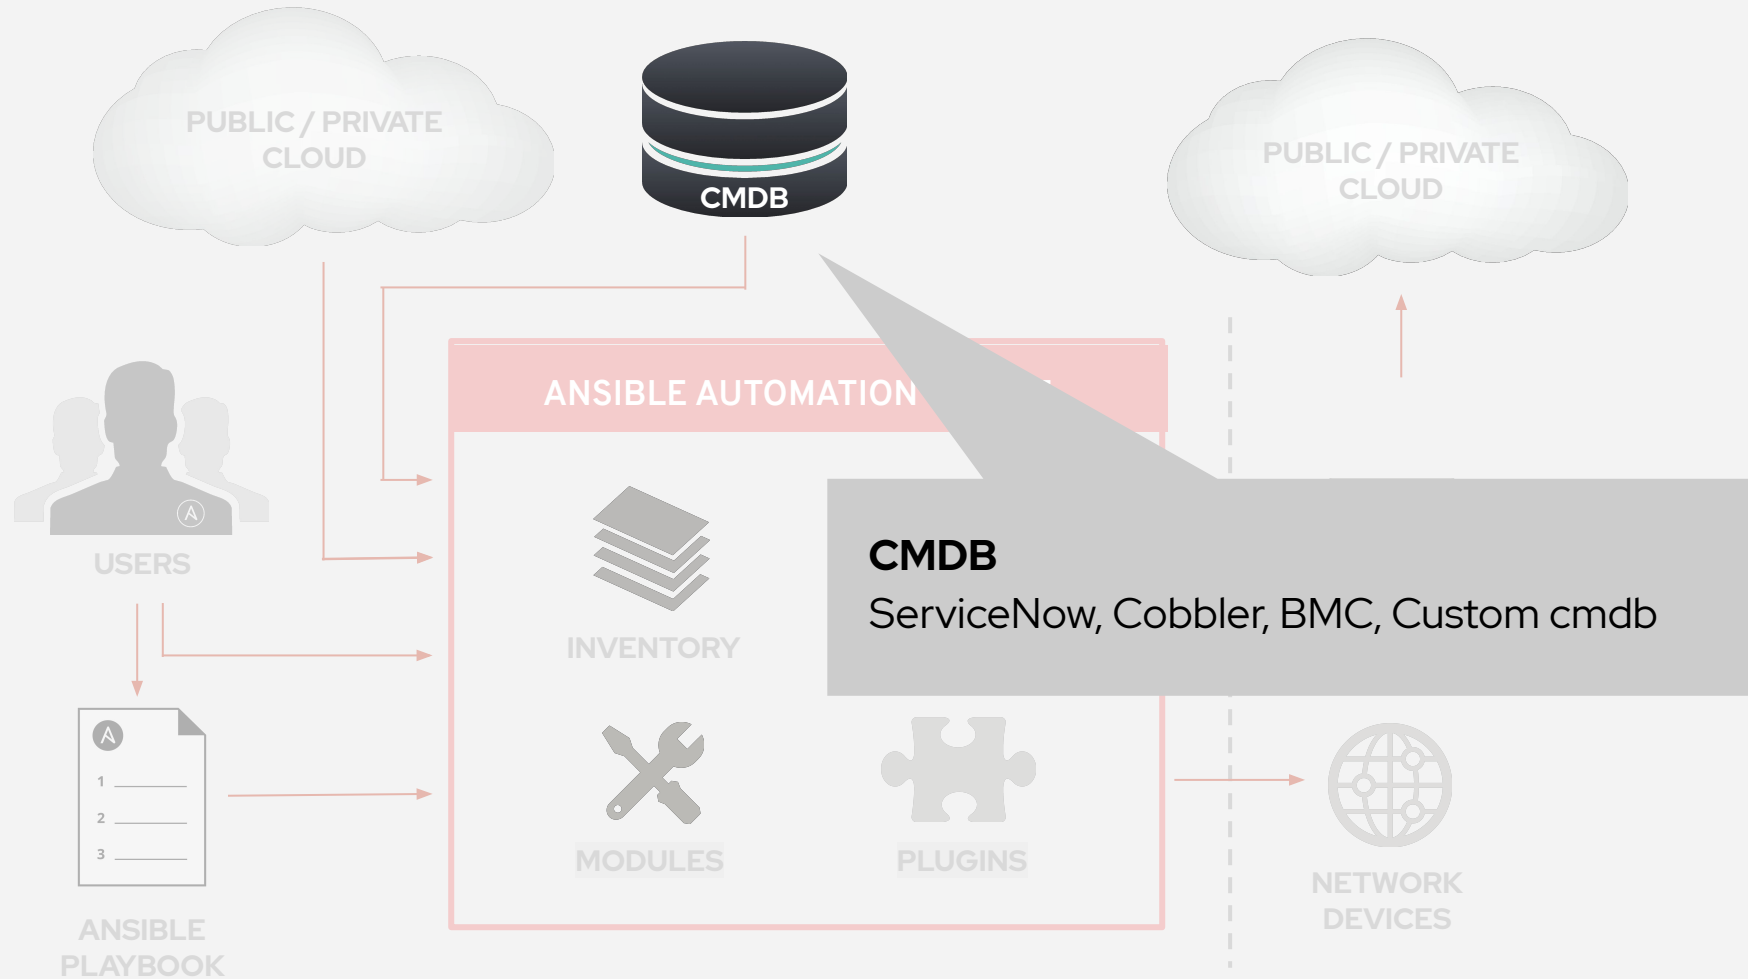


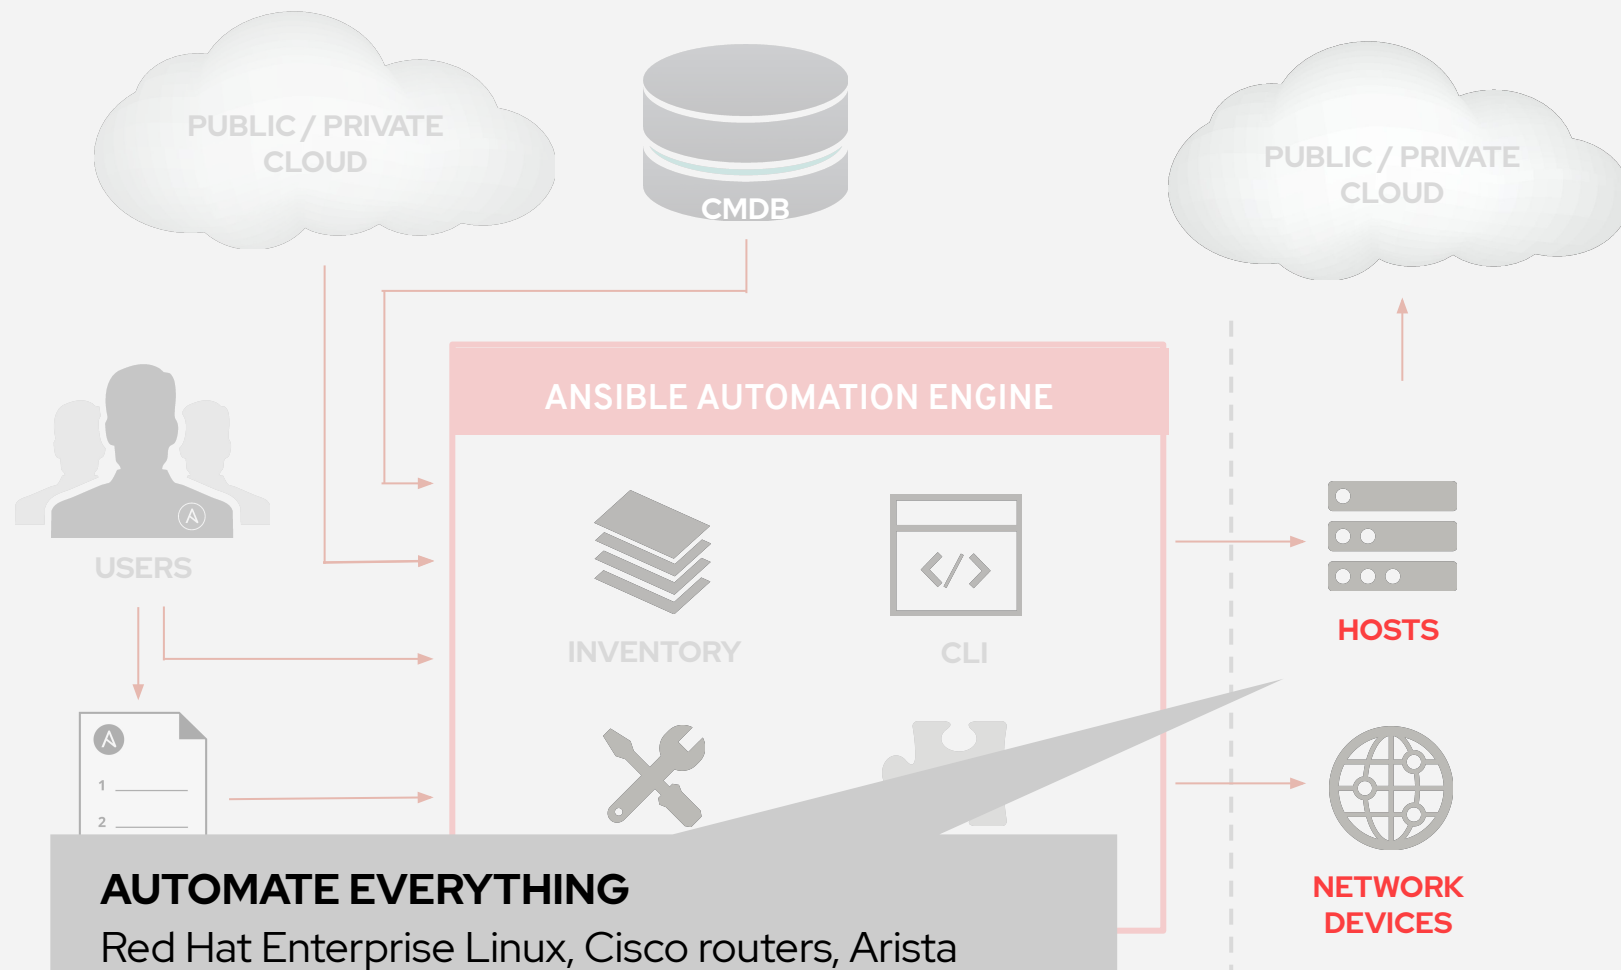
```
- name: latest index.html file is present
  template:
    src: files/index.html
    dest: /var/www/html/
```











AUTOMATE EVERYTHING

Red Hat Enterprise Linux, Cisco routers, Arista switches, Juniper routers, Windows hosts, Checkpoint firewalls, NetApp storage, F5 load balancers and more

Using Ansible

Ad-hoc commands

```
# check all my inventory hosts are ready to be  
# managed by Ansible  
$ ansible all -m ping
```

```
# run the uptime command on all hosts in the  
# web group  
$ ansible web -m command -a "uptime"
```

```
# collect and display the discovered for the  
# localhost  
$ ansible localhost -m setup
```

Ad-hoc example

Inventory

An inventory is a file containing:

- Hosts
- Groups
- Inventory-specific data (variables)
- Static or dynamic sources

Ansible Playbooks

```
- name: install and start apache
  hosts: web
  vars:
    http_port: 80
    max_clients: 200
  remote_user: root

  tasks:
    - name: install httpd
      yum: pkg=httpd state=latest
    - name: write the apache config file
      template: src=/srv/httpd.j2 dest=/etc/httpd.conf
    - name: start httpd
      service: name=httpd state=started
```

- name: install and start apache

hosts: web

vars:

http_port: 80

max_clients: 200

remote_user: root

tasks:

- name: install httpd

yum: pkg=httpd state=latest

- name: write the apache config file

template: src=/srv/httpd.j2 dest=/etc/httpd.conf

- name: start httpd

service: name=httpd state=started

- name: install and start apache

hosts: web

vars:

http_port: 80

max_clients: 200

remote_user: root

tasks:

- name: install httpd

yum: pkg=httpd state=latest

- name: write the apache config file

template: src=/srv/httpd.j2 dest=/etc/httpd.conf

- name: start httpd

service: name=httpd state=started

```
---
- name: install and start apache
  hosts: web
  vars:
    http_port: 80
    max_clients: 200
  remote_user: root

  tasks:
    - name: install httpd
      yum: pkg=httpd state=latest
    - name: write the apache config file
      template: src=/srv/httpd.j2 dest=/etc/httpd.conf
    - name: start httpd
      service: name=httpd state=started
```

```
---  
- name: install and start apache  
  hosts: web  
  vars:  
    http_port: 80  
    max_clients: 200  
    remote_user: root  
  
  tasks:  
    - name: install httpd  
      yum: pkg=httpd state=latest  
    - name: write the apache config file  
      template: src=/srv/httpd.j2 dest=/etc/httpd.conf  
    - name: start httpd  
      service: name=httpd state=started
```

```
---  
- name: install and start apache  
  hosts: web  
  vars:  
    http_port: 80  
    max_clients: 200  
    remote_user: root  
  
  tasks:  
    - name: install httpd  
      yum: pkg=httpd state=latest  
    - name: write the apache config file  
      template: src=/srv/httpd.j2 dest=/etc/httpd.conf  
    - name: start httpd  
      service: name=httpd state=started
```

tasks:

- name: add cache dir

file:

path: /opt/cache

state: directory

- name: install nginx

yum:

name: nginx

state: latest

notify: restart nginx

handlers:

- **name: restart nginx**

service:

name: nginx

state: restarted

Variables

Ansible can work with metadata from various sources and manage their context in the form of variables.

- Command line parameters
- Plays and tasks
- Files
- Inventory
- Discovered facts
- Roles

Tips/Best Practices

Simplicity

Simplicity

- hosts: web
tasks:
 - yum:
 - name: httpd
 - state: latest
- service:
 - name: httpd
 - state: started
 - enabled: yes

Simplicity

- hosts: web
name: install and start apache
tasks:
 - name: install apache packages
yum:
 - name: httpd
state: latest
 - name: start apache service
service:
 - name: httpd
state: started
enabled: yes

Naming example

Inventory

10.1.2.75

10.1.5.45

10.1.4.5

10.1.0.40

w14301.example.com

w17802.example.com

w19203.example.com

w19304.example.com

Inventory

db1 ansible_host=10.1.2.75

db2 ansible_host=10.1.5.45

db3 ansible_host=10.1.4.5

db4 ansible_host=10.1.0.40

web1 ansible_host=w14301.example.com



web2 ansible_host=w17802.example.com

web3 ansible_host=w19203.example.com

web4 ansible_host=w19203.example.com

Dynamic Inventories

- Stay in sync automatically
- Reduce human error



YAML Syntax

YAML and Syntax

- name: install telegraf
yum: name=telegraf-{{ telegraf_version }} state=present update_cache=yes
disable_gpg_check=yes enablerepo=telegraf
notify: restart telegraf
- name: configure telegraf
template: src=telegraf.conf.j2 dest=/etc/telegraf/telegraf.conf
- name: start telegraf
service: name=telegraf state=started enabled=yes

YAML and Syntax

- name: install telegraf
yum: >
 name=telegraf-{{ telegraf_version }}
 state=present
 update_cache=yes
 disable_gpg_check=yes
 enablerepo=telegraf
notify: restart telegraf
- name: configure telegraf
template: src=telegraf.conf.j2 dest=/etc/telegraf/telegraf.conf
- name: start telegraf
service: name=telegraf state=started enabled=yes

YAML and Syntax

- name: install telegraf
yum:
 - name: telegraf-{{ telegraf_version }}
 - state: present
 - update_cache: yes
 - disable_gpg_check: yes
 - enablerepo: telegrafnotify: restart telegraf
- name: configure telegraf
template:
 - src: telegraf.conf.j2
 - dest: /etc/telegraf/telegraf.confnotify: restart telegraf
- name: start telegraf
service:
 - name: telegraf
 - state: started
 - enabled: yes

```
ansible-playbook playbook.yml --syntax-check
```

Roles

Roles

- **Think about the full life-cycle of a service, microservice or container – not a whole stack or environment**
- **Keep provisioning separate from configuration and app deployment**
- **Roles are not classes or object or libraries – those are programming constructs**
- **Keep roles loosely-coupled – limit hard dependencies on other roles or external variables**

Variable Precedence

The order in which the same variable from different sources will override each other.

Variable Precedence

1. Extra vars
2. Include params
3. Role (and include_role) params
4. Set_facts / registered vars
5. Include_vars
6. Task vars (only for the task)
7. Block vars (only for tasks in the block)
8. Role vars
9. Play vars_files
10. Play vars_prompt
11. Play vars
12. Host facts / Cached set_facts
13. Playbook host_vars
14. Inventory host_vars
15. Inventory file/script host vars
16. Playbook group_vars
17. Inventory group_vars
18. Playbook group_vars/all
19. Inventory group_vars/all
20. Inventory file or script group vars
21. Role defaults
22. Command line values (e.g., -u user)

Things to Avoid

Things to Avoid

- Using command modules
 - Things like shell, raw, command etc.
- Complex tasks...at first
 - Start small
- Not using source control
 - But no really...

Ansible Content Collections

Collections Q and A

What are they?

- Collections are a distribution format for Ansible content that can include playbooks, roles, modules, and plugins. You can install and use collections through Ansible Galaxy and Automation Hub

How do I get them?

- `ansible-galaxy collection install namespace.collection -p /path`

Where can I get them?

- Today
 - Galaxy
 - Automation Hub

Collection Directory Structure

- **docs/**: local documentation for the collection
- **galaxy.yml**: source data for the MANIFEST.json that will be part of the collection package
- **playbooks/**: playbook snippets
 - **tasks/**: holds 'task list files' for include_tasks/import_tasks usage
- **plugins/**: all ansible plugins and modules go here, each in its own subdir
 - **modules/**: ansible modules
 - **lookups/**: lookup plugins
 - **filters/**: Jinja2 filter plugins
 - **connection/**: connection plugins required if not using default
- **roles/**: directory for ansible roles
- **tests/**: tests for the collection's content

Collections: Let's Go!

1. Init collection: `ansible-galaxy collection init foo.bar`
2. Sanity testing: `ansible-test sanity`
3. Unit tests: `ansible-test units`
4. Integration tests: `ansible-test integration`
5. Build the collection: `ansible-galaxy collection build`
6. Publish the collection: `ansible-galaxy collection publish`
7. Install the collection: `ansible-galaxy collection install
foo.bar`

Thank you

Questions?

 linkedin.com/company/red-hat

 youtube.com/user/RedHatVideos

 facebook.com/redhatinc

 twitter.com/RedHat

Resource Link Index

https://docs.ansible.com/ansible/latest/user_guide/playbooks_variables.html#variable-precedence-where-should-i-put-a-variable
https://docs.ansible.com/ansible/latest/user_guide/playbooks_variables.html#using-variables
https://docs.ansible.com/ansible/latest/user_guide/playbooks_intro.html
https://docs.ansible.com/ansible/latest/installation_guide/intro_installation.html
https://docs.ansible.com/ansible/latest/user_guide/intro_getting_started.html#getting-started
https://docs.ansible.com/ansible/latest/user_guide/intro_adhoc.html
https://docs.ansible.com/ansible/latest/user_guide/intro_inventory.html
<https://docs.ansible.com/ansible/latest/index.html>
https://docs.ansible.com/ansible/latest/user_guide/playbooks_reuse_roles.html
https://docs.ansible.com/ansible/latest/user_guide/intro_dynamic_inventory.html
<https://docs.ansible.com/ansible-lint/>
<https://github.com/ansible/ansible>
<https://github.com/ansible/ansible-lint>
<https://ansible.github.io/workshops/>
<https://www.ansible.com/resources/ebooks/get-started-with-red-hat-ansible-tower>
https://docs.ansible.com/ansible/latest/user_guide/collections_using.html
https://docs.ansible.com/ansible/latest/dev_guide/developing_collections.html